



# MongoDB and Cloud Storage

Frank Walsh, Diarmuid O'Connor

# Agenda

- Cloud Databases
- MongoDB
  - Querying
  - Integrating with Node.js
  - The Contacts API implementation

# Databases in Enterprise Apps

- Most data driven enterprise applications need a database
- In traditional enterprise applications, this requires
  - Backups
  - Fail over
  - Maintenance
  - Capacity provisioning
- Usually handled by a Database Administrator.

# Databases in the Cloud

- For some apps, a traditional database may not be the best fit
  - Does the app require transactional integrity
  - Do you need db schema definition
  - Do you know your scaling requirements, particularly if it's a web app
- One approach is to use the **Cloud** for you DB
  - Designed for scale
  - Can be outsourced so you don't have to deal with infrastructure requirements.

# Cloud DB Advantages

- Removes Management costs
- Inherently scalable
- Latency is predictable and constant
- No need to define schemas etc.
- Lots of Cloud DB offerings out there
  - SQL based
  - NoSQL based
- If organisation policy/standards do not allow outsourcing:
  - Can host yourself, most NoSQL DBs are free.

# Cloud Database Practices

- Drop Consistency
  - this makes distributed systems much easier to build
- Drop SQL and the relational model
  - simpler structures are easier to distribute:
    - key/value pairs
    - **structured documents**
    - **pseudo-tables**
    - tend to be schema-free, accepting data as-is
- Offer HTTP interfaces using XML or **JSON**
- Use in-memory storage aggressively

# Designing Distributed Data

- App data is not homogeneous
  - some kinds of data will be much larger
- consider using different databases for different requirements:
- user details,billing - needs consistency
  - require traditional database
- user data,content - needs partition tolerance
  - replicate to keep safe
- analytics,sessions - needs availability
  - "eventually consistent" is good enough

**MONGODB**



# Introduction

- Document-oriented database
  - but closer to traditional SQL databases than others
- Uses JSON natively - perfect fit for Node.js
- Query language with many SQL features
  - Uses JSON too, and has an easy learning curve
- Inbuilt sharding support means you can scale
  - Aggressively uses memory for high speed
- be careful: default configuration does not sync to disk
- Commercial support - 10gen.com product
  - cloud hosting providers - e.g. mongoLab.com
- Community support - popular choice

# Mongo Terminology

- Each database contains a set of "Collections"
- Collections are analogous to SQL tables
- Collections contain a set of JSON documents
  - there is no schema (in the DB)
- the documents can all be different
  - means you have rapid development
  - adding a property is easy - just starting using in your code
- makes deployment easier and faster
  - roll-back and roll-forward are safe - unused properties are just ignored
- Collections can be indexed and queries
- Operations on individual documents are atomic

# Getting Started (locally)

- For complete MongoDB installation instructions, see [the manual](#).
- Starting MongoDB:  
`mongod`
- This starts the process.
- Can add other parameters, for instance location of data.

# Mongo Shell

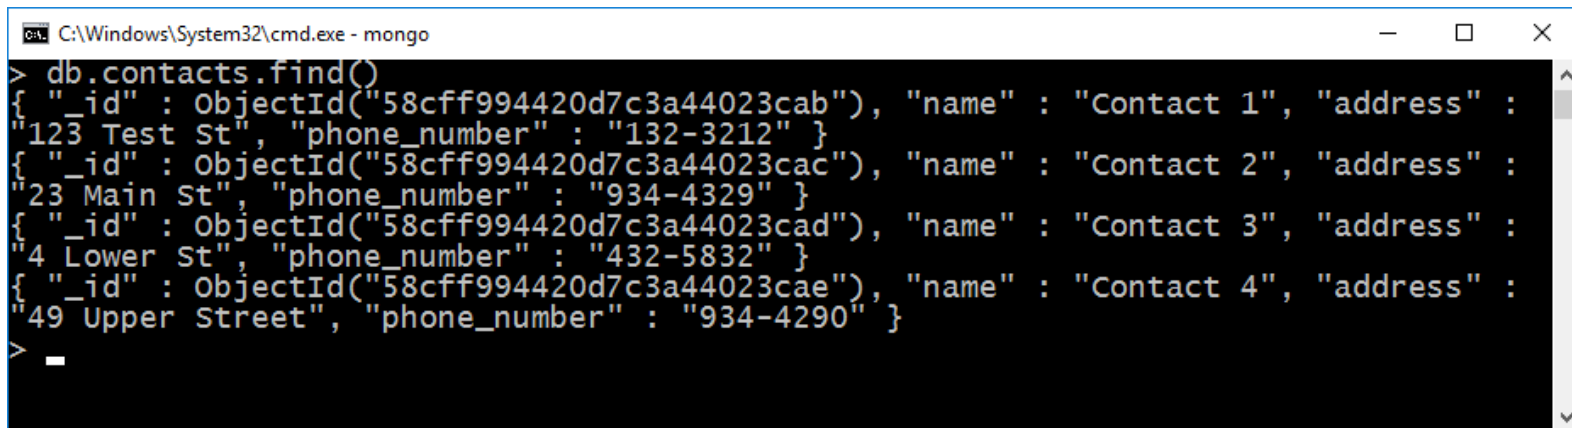
- Interactive JavaScript interface to MongoDB.
- Query/update data and perform administrative operations.

```
C:\repos\webervicesdev-2017>mongo
MongoDB shell version v3.4.2
connecting to: mongodb://127.0.0.1:27017
MongoDB server version: 3.4.2
Server has startup warnings:
2017-03-20T13:49:38.768+0000 I CONTROL [initandlisten]
2017-03-20T13:49:38.769+0000 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database.
2017-03-20T13:49:38.773+0000 I CONTROL [initandlisten] **      Read and write access to data and configuration is u
nrestricted.
2017-03-20T13:49:38.775+0000 I CONTROL [initandlisten]
>
```

- By default, Mongo shell will attempt to connect to the MongoDB instance running on the localhost interface on port 27017.

# The MongoDB Query Language

- MongoDB provides a JavaScript API and JSON-based query language
- Use the MongoDB shell to execute queries
  - similar to using MySQL console
- Example: list of contacts

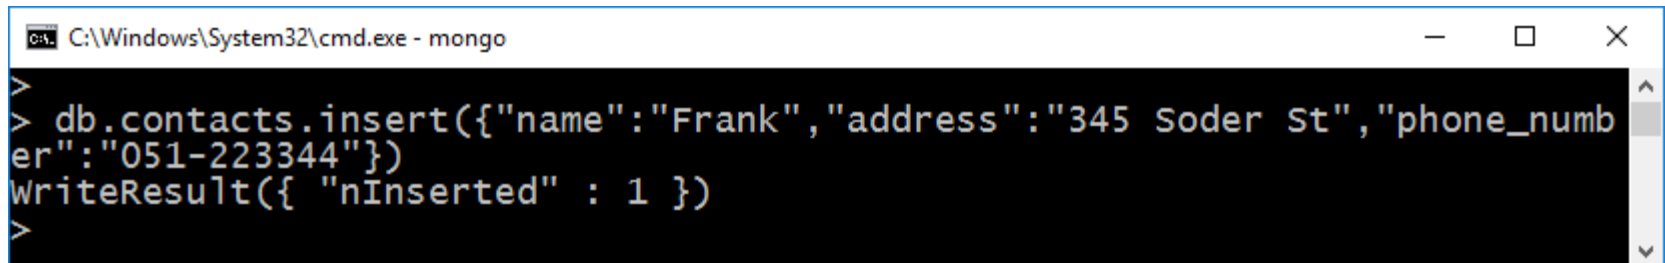


```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find()
{ "_id" : ObjectId("58cff994420d7c3a44023cab"), "name" : "Contact 1", "address" :
123 Test St", "phone_number" : "132-3212" }
{ "_id" : ObjectId("58cff994420d7c3a44023cac"), "name" : "Contact 2", "address" :
23 Main St", "phone_number" : "934-4329" }
{ "_id" : ObjectId("58cff994420d7c3a44023cad"), "name" : "Contact 3", "address" :
4 Lower St", "phone_number" : "432-5832" }
{ "_id" : ObjectId("58cff994420d7c3a44023cae"), "name" : "Contact 4", "address" :
49 Upper Street", "phone_number" : "934-4290" }
>
```

- db - current database
- contacts - the contacts collection
- .find() - collection API method (corresponds to collection URL in last lecture...)
- The Result Set is a list of JavaScript objects, representing matched documents

# MongoDB: Inserts

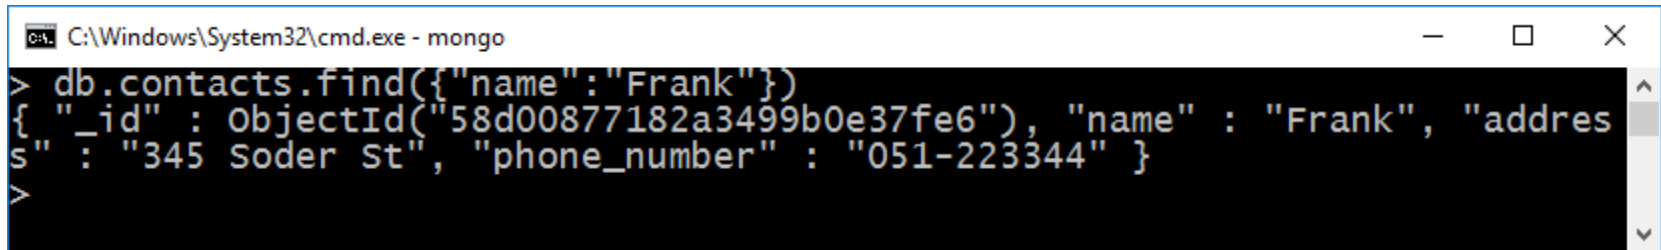
- Collections do not need to be created explicitly
  - just insert a document
- MongoDB automatically assigns a 12 byte unique identifier to any document
  - the **\_id** property
  - Stored internally as binary

A screenshot of a Windows command prompt window titled "C:\Windows\System32\cmd.exe - mongo". The window has a black background with white text. The command prompt shows a MongoDB shell session with the following text: > db.contacts.insert({"name":"Frank","address":"345 Soder St","phone\_number":"051-223344"}) followed by the output WriteResult({ "nInserted" : 1 }). The prompt ends with >.

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.insert({"name":"Frank","address":"345 Soder St","phone_number":"051-223344"})
WriteResult({ "nInserted" : 1 })
>
```

# MongoDB:Queries

- Documents are retrieved by specifying a set of conditions to match against
- simplest case : query-by-example
- provide a subset of properties that must match



```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":"Frank"})
{"_id" : ObjectId("58d00877182a3499b0e37fe6"), "name" : "Frank", "address" : "345 Soder St", "phone_number" : "051-223344" }
```

- More complex queries use a convention of embedded meta- properties to specify conditions these are signified with a \$ prefix Example:{name:{\$exists:true}}  
returns documents that have a name property

# MongoDB: Queries

- Common meta-properties used to query data are:
  - **\$gt, \$gte, \$lt, \$lte** meaning >, >=, <, <=

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.insert({"name":"Mary","age":21,"address":"345 Soder St","phone_number":"051-223344"})
WriteResult({ "nInserted" : 1 })
> db.contacts.insert({"name":"Jane","age":31,"address":"345 Keel St","phone_number":"051-445566"})
WriteResult({ "nInserted" : 1 })
> db.contacts.find({"age":{"$gte":21,$lt:31}})
{"_id" : ObjectId("58d00909182a3499b0e37fe7"), "name" : "Mary", "age" : 21, "address" : "345 Soder St", "phone_number" : "051-223344" }
```

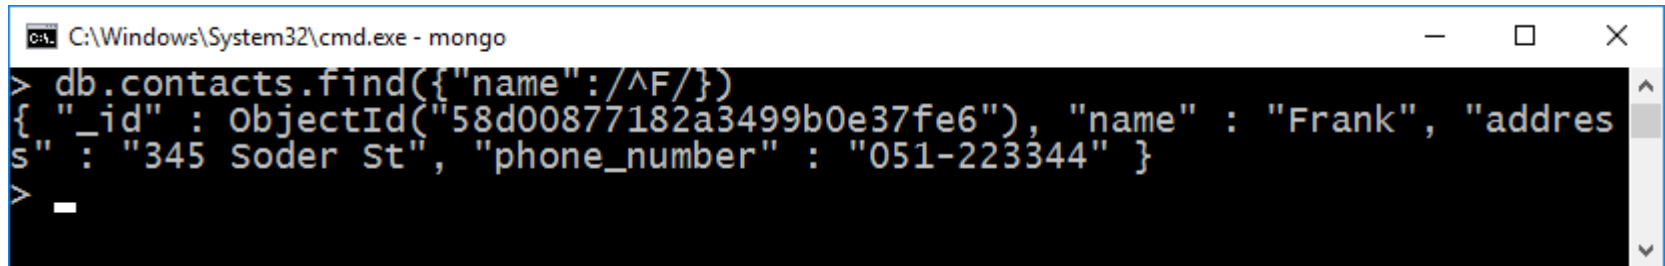
- **\$or, \$in, \$nin**

```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":{"$in":["Jane","Frank"]}})
{"_id" : ObjectId("58d00877182a3499b0e37fe6"), "name" : "Frank", "address" : "345 Soder St", "phone_number" : "051-223344" }
{"_id" : ObjectId("58d00939182a3499b0e37fe8"), "name" : "Jane", "age" : 31, "address" : "345 Keel St", "phone_number" : "051-445566" }
```



# MongoDB: Queries

- **regular expressions** {word: /th^/i }



```
C:\Windows\System32\cmd.exe - mongo
> db.contacts.find({"name":/^F/})
{"_id" : ObjectId("58d00877182a3499b0e37fe6"), "name" : "Frank", "address" : "345 Soder St", "phone_number" : "051-223344" }
>
```

- `db.contacts.find().limit(5)`
  - limits the number of documents in the result set.
- `db.contacts.find().skip( 5 )`
  - Set the Starting Point of the Result Set

# MongoDB:Updates

- Documents are updated by providing:
  - a query to select the relevant subset of documents,
  - an update specification, which is either:
    - a complete replacement document
    - meta-properties that modify specific document properties
- example:  
**\$set** changes specific properties  
Example:complete replacement:  
> db.city.insert( {name:'dublin'} )  
> db.city.update( {name:'dublin'}, {name:'Dublin',county:'Dublin'} )
- Example:modify specific properties:  
> db.city.insert( {name:'Cork',county:'cork'} )  
> db.city.update( {name:'Cork'}, {\$set:{county:'Cork'}} )
- See <http://www.mongodb.org/display/DOCS/Updating> for more

# MongoDB:Update Properties

- Common meta-properties used with the update command are:
  - **\$set** - sets specified properties, but leaves others alone  
\$set:{name:'New Name'}
- **\$unset** - deletes specified properties  
\$unset:{name:1}
- **\$inc** - increments a numeric property  
inc:{ upvotes: 2 }  
adds 2 to the counter property, or if it does not exist, sets it to 2
- **\$push, \$pop** - add to or remove values from, an array
  - \$push: { comments: {who:..., msg:...} }
  - \$pop: {comments: -1 }

# MongoDB:Upserts

- The MongoDB update command can optionally insert a document if it is not found. This is known as an 'upsert'
- This is useful when starting counters as it avoids corrupting the count when two independent updates try to initialize the counter.

```
db.counters.update( {name:'foo'}, {$inc:{value:1}},  
true)
```

- The first update will create the counter:  
{name:'foo', value:1}
- The second update will increment the counter:  
{name:'foo', value:2}

# **MONGO DB NODE.JS DRIVER**

# MongoDB Node.jsDriver

- To connect Node.js to MongoDB you need a *database driver*
- A Node.js module that communicates over the wire to MongoDB, and presents an API over the database.
- The one we'll look at available with NPM:

<http://docs.mongodb.org/ecosystem/drivers/node-js/>

```
npm install mongodb
```

```
var mongodb = require('mongodb')
```

# node-mongodb-native

- Provides low-level interface to MongoDB and replicates the MongoDB Console
  - suffers badly from callbacks !
- Used by higher level MongoDB modules:
  - mongoose:Object Relational Mapper
  - mongoskin:simplerAPI to reduce callbacks
- Sufficient for smaller apps

# Connecting to MongoDB using Node

```
const mongo = require('mongodb');
const MongoClient = mongo.MongoClient;
const mongoDb;
const url = `mongodb://localhost:27017/myproject`
MongoClient.connect(url, (err, db)=> {    if(!err) {
    console.log("We are connected");
    mongoDb = db; }
    else
        {console.log("Unable to connect to
the db");
    }
}
);
```



# Inserting a Document

The following function that will insert a document:

```
if (mongoDb) {  
    const collection = mongoDb.collection('contacts');  
    collection.insert(contact, {w:1}, (err, result)=> {  
        if (err) {  
            console.log({'error':'An error has occurred'});  
        } else {  
            console.log('Success: ' + JSON.stringify(result[0]));  
        }  
    });  
}  
else  
{  
    console.log('No database object!');  
}
```

- The insert function returns a result object that contains several fields that might be useful.
  - **result** Contains the result document from MongoDB
  - **ops** Contains the documents inserted with added **\*\*\_id\*\*** fields
  - **connection** Contains the connection used to perform the insert

# Updating a Document

- The following simple document update by adding a new field **b** to the document that has the field **a** set to **2**.

```
// Get the documents collection
const collection = db.collection('documents');
// Update document where a is 2, set b equal to 1
collection.update({ a : 2 }, { $set: { b : 1 } }, (err,
result)=>{
    console.log("Updated the document with
the field a equal to 2");
});
```

- The method will update the first document where the field **a** is equal to **2** by adding a new field **b** to the document set to **1**.

# Deleting a Document

- This will remove the first document where the field **a** equals to **3**.

```
var collection = db.collection('documents');
// Insert some documents
collection.remove({ a : 3 }, (err, result)=>{
    If (!err){
        console.log("Removed the document with the
field a          equal to 3");}
    }
    else{
        console.log("Did not remove document");}

})
}
```

**MONGOOSE**

# Mongoose Overview

- Mongoose is a object-document model module in Node.js for MongoDB
  - Wraps the functionality of the native MongoDB driver
  - Exposes models to control the records in a doc
  - Supports validation on save
  - Extends the native queries

# Installing Mongoose

- Run the following from the CMD/Terminal
  - \$ npm install -save mongoose
- In node
  - Load the module
    - import mongoose from ('mongoose');
- Connect to the database
  - mongoose.connect(mongoDbPath);

# Mongoose Schemas and Models

- Mongoose supports models
  - i.e. fixed types of documents
  - Needs a mongoose.Schema
  - Each of the properties must have a type
    - Number, String, Boolean, array, object

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5      name: String,
6      address: String,
7      age: Number,
8      email: String,
9      updated: Date
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```

# Mongoose Schema - Validation

- Can define validation constraints on properties :

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5    name: String,
6    address: String,
7    age: { type: Number, min: 0, max: 120 },
8    email: String,
9    updated: { type: Date, default: Date.now }
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```



# Mongoose Schemas

## Arrays

Comments property is  
an Array of  
CommentSchemas

```
1 |const mongoose = require('mongoose'),
2 |Schema = mongoose.Schema;
3
4 |const CommentSchema = new Schema({
5 |  body: {type: String, required:true},
6 |  author: {type: String, required:true},
7 |  upvotes:Number
8 |});
9
10 |const PostSchema = new Schema({
11 |  title: {type: String, required:true},
12 |  link: {type: String, optional:true},
13 |  username: {type: String, required:true},
14 |  comments: [CommentSchema],
15 |  upvotes: { type: Number, min: 0, max: 100 }
16 |});
17
18 |export default mongoose.model('posts', PostSchema);
```

# Mongoose Custom Validation

- Developers can define custom validation on their properties (e.g. validate length of comment when trying to save)

```
18 CommentSchema.path('body').validate((v)=>{
19     if (v.length>40 || v.length < 5){
20         return false
21     }
22     return true
23 })
24
```

# Data Manipulation Mongoose

- Mongoose supports all the CRUD operations:
  - Create → `model.create(callback)`
  - Read → `Model.find().exec(callback)`
  - Update → `modelObj.update(props, callback)` → `Model.update(condition, props, cb)`
  - Remove → `modelObj.remove(callback)` → `Model.remove(condition, props, cb)`

# Create Contact with Mongoose

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5    name: String,
6    address: String,
7    age: { type: Number, min: 0, max: 120 },
8    email: String,
9    updated: { type: Date, default: Date.now }
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```

```
1  router.post('/', (req, res) => {
2    let newContact = req.body;
3    if (newContact){
4      Contact.create(newContact, (err, contact) => {
5        if(err) { return handleError(res, err); }
6        return res.status(201).send({contact});
7      });
8    }else{
9      return handleError(res, err);
10   }
11 });
```

# Update Contact with Mongoose

```
1  const mongoose = require('mongoose'),
2  Schema = mongoose.Schema;
3
4  ▼ const ContactSchema = new Schema({
5    name: String,
6    address: String,
7    age: { type: Number, min: 0, max: 120 },
8    email: String,
9    updated: { type: Date, default: Date.now }
10  });
11
12  const ContactModel = mongoose.model('contacts', ContactSchema);
```

```
1  router.put('/:id', (req, res) => {
2    let key = req.params.id;
3    let updateContact = req.body;
4
5    if(updateContact._id) {delete updateContact._id;}
6    ContactModel.findById(req.params.id, (err, contact) => {
7      if (err) { return handleError(res, err); }
8      if(!contact) { return res.send(404); }
9      const updated = _.merge(contact, req.body);
10     updated.save((err) => {
11       if (err) { return handleError(res, err); }
12       return res.send(contact);
13     });
14   });
15 });
16
```

# Mongoose Queries

- Mongoose provides a more expressive version of the native MongoDB
  - Instead of:  
`{ $or: [{ conditionOne: true }, { conditionTwo: true } ] }`
  - Do:  
`.where({ conditionOne: true }).or({ conditionTwo: true })`

# Mongoose Queries

- Mongoose supports many queries:
  - For equality/non-equality
  - Selection of some properties
  - Sorting
  - Limit & skip
- All queries are executed over the object returned by `Model.find*()`
  - `Model.findOne()` returns a single document, the first match
  - `Model.find()` returns all
  - `Model.findById()` queries on the `_id` field.

```
1
2   let id = "58cff994420d7c3a44023cb0"
3   ContactModel.findById(id,(err, contact) => {
4       if (err) { return handleError(res, err); }
5       console.log(contact)
6   });
7
```

# Mongoose Queries

- Can build complex queries and execute them later

```
1  const query = ContactModel.where('age').gt(17).lt(66)
2    .where('county').in(['Waterford', 'Wexford', 'Kilkenny']);
3
4  query.exec((err, contacts) => {...})
5
6
```

- The above finds all contacts where age >17 and <66 and living in either Waterford, Kilkenny or Wexford



# Mongoose Sub-Docs

- Ex: Hacker News – Adding a comment to a post.

```
1  //add comment
2  router.post('/:id/comments', (req, res) => {
3      const id = req.params.id;
4      const comment = req.body;
5      PostModel.findById(id, (err, post)=>{
6          if(err) { return handleError(res, err); }
7          post.comments.push(comment);
8          post.save(err => {
9              if (err) {return handleError(res, err);}
10             return res.status(201).send({post});
11         });
12     });
13 });
```

# Mongoose Sub-Docs

- Updating a Sub-Document(e.g. incrementing the upvotes for a comment)

```
1  router.post('/:postId/comments/:commentId/upvotes', (req, res) => {
2      const commentId = req.params.commentId;
3      const postId = req.params.postId;
4      Post.findById( postId, (err, post)=>{
5          if(err) { return handleError(res, err); }
6          post.comments.id(commentId).upvotes++;
7          post.save(err => {
8              if (err) {return handleError(res, err);}
9              return res.status(201).send({post});
10         });
11     });
12 });
```

Each subdocument is assigned its own `_id` from MongoDB. This is a special method to access sub documents

# Mongo Sub docs

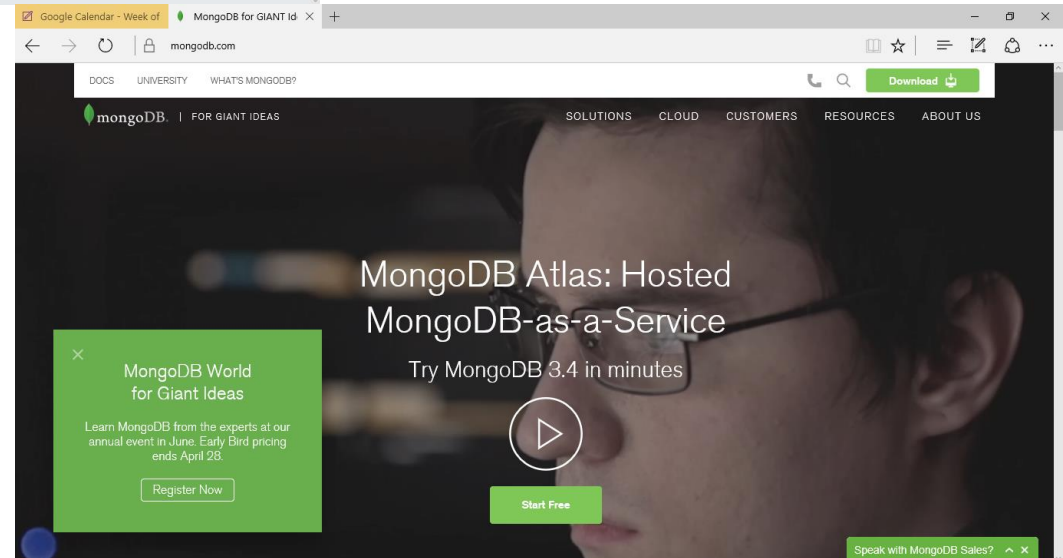
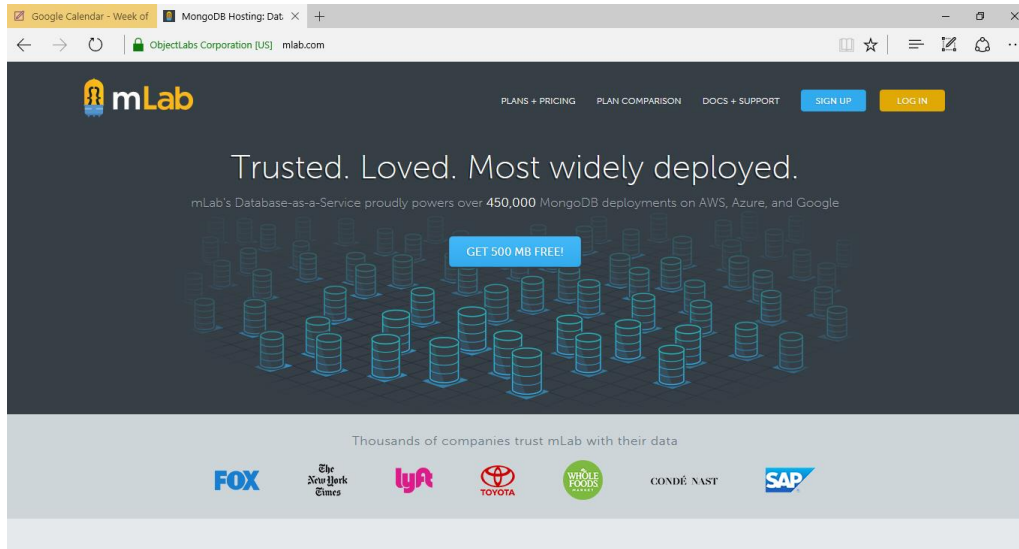
- Removing a sub document

```
218 Post.findById('5510117c1a9f03cd1ed38a6c', function (err,p){
219     p.comments.id('551020f317bf07692231caed').remove()
220     p.save (function (err) {
221         if (err) {
222             console.log(p)
223         } else {
224             console.log('comment removed')
225         }
226         process.exit()
227     })
228 }
```

# MongoDB as a Service


- Best practice for initial development is to host MongoDB process on your development machine
- In production environments, Mongo will be hosted:
  - on it's own instance or
  - provisioned as a service

# MongoDB as a Service



# MongoDB as a Service

- Most providers allow free access to their
- Provide user credentials wrapped in a URL
- All you need to do is update your config with the relevant URL
- Careful to ignore credentials when pushing to github/public repo



[Home](#)  
Database: contacts\_db

To connect using the mongo shell:

```
% mongo ds039311.mlab.com:39311/contacts_db -u <dbuser> -p <dbpassword>
```

To connect using a driver via the standard MongoDB URI ([what's this?](#)):

```
|mongodb://<dbuser>:<dbpassword>@ds039311.mlab.com:39311/contacts_db
```